



UNIMORE
UNIVERSITÀ DEGLI STUDI DI
MODENA E REGGIO EMILIA

Dipartimento di Scienze Fisiche,
Informatiche e Matematiche

14. Compilatori e Parallel Programming Models

Linguaggi e Compilatori [I215-011]

Corso di Laurea in INFORMATICA
(D.M.270/04) [16-262]
Anno accademico 2022/2023


Prof. Andrea Marongiu
andrea.marongiu@unimore.it

Programming model: OpenMP


- De-facto standard for the **shared memory** programming model
- A collection of **compiler directives**, **library routines** and **environment variables**
- Easy to specify parallel execution within a **serial code**
- Targets several types of parallelism (data, task, heterogeneous)
- Popular in high-end embedded
- Requires **special support** in the compiler
- Generates calls to **threading libraries** (e.g. pthreads)

Fork/Join Parallelism

Sequential program



Parallel program



- Initially only master thread is active
- Master thread executes sequential code
- Fork: Master thread creates or awakens additional threads to execute parallel code
- Join: At the end of parallel code created threads are suspended upon **barrier** synchronization

Pragmas

- **Pragma**: a compiler directive in C or C++
- Stands for “pragmatic information”
- A way for the programmer to communicate with the compiler
- Compiler free to ignore pragmas: original sequential semantic is not altered
- Syntax:

#pragma omp *<rest of pragma>*

Components of OpenMP

a subset of the directives

Directives

- ❖ Parallel regions
 - *#pragma omp parallel*
- ❖ Work sharing
 - *#pragma omp for*
 - *#pragma omp sections*
- ❖ Synchronization
 - *#pragma omp barrier*
 - *#pragma omp critical*
 - *#pragma omp atomic*

Runtime Library

Clauses

- ❖ Data scope attributes
 - *private*
 - *shared*
 - *reduction*
- ❖ Loop scheduling
 - *static*
 - *dynamic*

- ❖ Thread Forking/Joining
 - *omp_parallel_start()*
 - *omp_parallel_end()*
- ❖ Loop scheduling
- ❖ Thread IDs
 - *omp_get_thread_num()*
 - *omp_get_num_threads()*

Outlining parallelism

The **parallel** directive

```
void ex1 ()
{
  <bb 2> [0.00%]:
  __builtin_GOMP_parallel (
    ex1._omp_fn.0, 0B, 0, 0);
}
```

```
ex1._omp_fn.0 (void * .omp_data_i)
{
  <bb 3> [0.00%]:
  printf ("Hello World!\n");
  return;
}
```

[hello.c.012t.ompexp](#)

Code originally contained within the scope of the pragma is outlined to a new function within the compiler

**A sequential program..
..is easily parallelized**

- ☐ Fundamental construct to outline parallel computation within a sequential program
- ☐ Code within its scope is **replicated** among threads
- ☐ Defers implementation of parallel execution to the runtime (machine-specific, e.g. [pthread_create](#))

```
void ex1 ()
{
  #pragma omp parallel
  {
    printf ("\nHello world!");
  }
}
```

Outlining parallelism

The **parallel** directive

```
void ex1 ()
{
    <bb 2> [0.00%]:
    __builtin_GOMP_parallel (
        ex1._omp_fn.0, 0B, 0, 0);
}

ex1._omp_fn.0 (void * .omp_data_i)
{
    <bb 3> [0.00%]:
    printf ("Hello World!\n");
    return;
}
hello.c.012t.ompexp
```

The `#pragma` construct in the **main** function is replaced with function calls to the runtime library

```
void ex1 ()
{
    #pragma omp parallel
    {
        printf ("\nHello world!");
    }
}
```

NOTE: The **master** thread is the only thread that executes this code

#pragma omp parallel

Within the runtime we first call a function to fork new threads,

```
void ex1 ()
{
    <bb 2> [0.00%]:
    __builtin_GOMP_parallel (
        ex1.omp_fn.0, 0B, 0, 0);
}

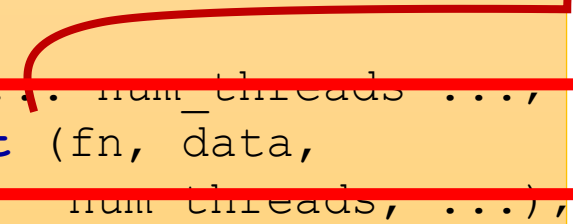
ex1.omp_fn.0 (void * .omp_data_i)
{
    <bb 3> [0.00%]:
    printf ("Hello World!\n");
    return;
}
```

[hello.c.012t.ompexp](#)

```
void ex1 ()
{
    #pragma omp parallel
    {
        printf ("\nHello world!");
    }
}
```

```
void
GOMP_parallel (void (*fn) (void *),
               void *data,
               unsigned num_threads,
               unsigned int flags)
{
    num_threads = ... num_threads ...;
    gomp_team_start (fn, data,
                    num_threads, ...);
    fn (data);
    ialias_call (GOMP_parallel_end) ();
}
```

GOMP_parallel_start



#pragma omp parallel

Within the runtime we first call a function to fork new threads, **passing a pointer to the outlined function**

```
void ex1 ()
{
    <bb 2> [0.00%]:
    __builtin_GOMP_parallel (
        ex1._omp_fn.0, 0B, 0, 0);
}
```

GOMP_parallel_start

```
ex1._omp_fn.0 (void * .omp_data_i)
{
    <bb 3> [0.00%]:
    printf ("Hello World!\n");
    return;
}
```

[hello.c.012t.ompexp](#)

```
void
GOMP_parallel (void (*fn) (void *),
               void *data,
               unsigned num_threads,
               unsigned int flags)
{
    num_threads = ... num_threads ...;
    gomp_team_start (fn, data,
                    num_threads, ...);
    fn (data);
    ialias_call (GOMP_parallel_end) ();
}
```

```
void ex1 ()
{
    #pragma omp parallel
    {
        printf ("\nHello world!");
    }
}
```

...and **data**, which we will be looking at in a moment...

#pragma omp parallel

```
void ex1 ()
{
    <bb 2> [0.00%]:
    __builtin_GOMP_parallel (
        ex1._omp_fn.0, 0B, 0, 0);
}

ex1._omp_fn.0 (void * .omp_data_i)
{
    <bb 3> [0.00%]:
    printf ("Hello World!\n");
    return;
}
```

[hello.c.012t.ompexp](#)

```
void ex1 ()
{
    #pragma omp parallel
    {
        printf ("\nHello world!");
    }
}
```

Then the master itself calls
the parallel function

```
void
GOMP_parallel (void (*fn) (void *),
               void *data,
               unsigned num_threads,
               unsigned int flags)
{
    num_threads = ... num_threads ...;
    gomp_team_start (fn, data,
                     num_threads, ...);
    fn (data);
    _alias_call (GOMP_parallel_end) ();
}
```

REMEMBER: The master
thread is the only thread that
executes this library code

#pragma omp parallel

```
void ex1 ()
{
    <bb 2> [0.00%]:
    __builtin_GOMP_parallel (
        ex1._omp_fn.0, 0B, 0, 0);
}

ex1._omp_fn.0 (void * .omp_data_i)
{
    <bb 3> [0.00%]:
    printf ("Hello World!\n");
    return;
}
```

[hello.c.012t.ompexp](#)

```
void ex1 ()
{
    #pragma omp parallel
    {
        printf ("\nHello world!");
    }
}
```

Finally we call the runtime to
synchronize threads with a
barrier and suspend them

```
void
GOMP_parallel (void (*fn) (void *),
               void *data,
               unsigned num_threads,
               unsigned int flags)
{
    num_threads = ... num_threads ...;
    gomp_team_start (fn, data,
                     num_threads, ...);
    fn (data);
    ialias_call (GOMP_parallel_end) ();
}
```

REMEMBER: The master
thread is the only thread that
executes this library code

GOMP: The OpenMP runtime library

- Where does this library code live?
- GCC implementation → **libgomp**
<https://github.com/gcc-mirror/gcc/tree/master/libgomp>

```
hero-vm@ubuntu:~/hero-sdk/hero-gcc-toolchain/src/riscv-gcc/libgomp$ ls
acinclude.m4      hashtable.h       Makefile.in      parallel.c
aclocal.m4        icv.c            oacc-async.c     plugin
affinity.c        icv-device.c     oacc-cuda.c      priority_queue.c
alloc.c           iter.c           oacc-host.c      priority_queue.h
atomic.c          iter_uhl.c       oacc-init.c      sections.c
barrier.c         libgomp_f.h.in   oacc-int.h       single.c
ChangeLog         libgomp_g.h      oacc-mem.c       splay-tree.c
ChangeLog.graphite libgomp.h         oacc-parallel.c  splay-tree.h
config            libgomp.map      oacc-plugin.c    target.c
config.h.in       libgomp-plugin.c oacc-plugin.h    task.c
configure         libgomp-plugin.h omp.h.in          taskloop.c
configure.ac       libgomp.spec.in  omp_lib.f90.in   team.c
configure.tgt     libgomp.texi     omp_lib.h.in     testsuite
critical.c        lock.c           openacc.f90      work.c
env.c             loop.c           openacc.h
error.c           loop_uhl.c       openacc_lib.h
fortran.c         Makefile.am      ordered.c
```

Easy to spot at least one file per #pragma construct

GOMP: The OpenMP runtime library

- Where does this library code live?
- GCC implementation → **libgomp**
<https://github.com/gcc-mirror/gcc/tree/master/libgomp>

```
hero-vm@ubuntu:~/hero-sdk/hero-gcc-toolchain/src/riscv-gcc/libgomp$ ls
acinclude.m4      hashtable.h       Makefile.in      parallel.c
aclocal.m4        icv.c            oacc-async.c     plugin
affinity.c        icv-device.c     oacc-cuda.c      priority_queue.c
alloc.c           iter.c           oacc-host.c      priority_queue.h
atomic.c          iter_ull.c       oacc-init.c      sections.c
barrier.c         libgomp_f.h.in   oacc-int.h       single.c
ChangeLog         libgomp_g.h      oacc-mem.c       splay-tree.c
ChangeLog.graphite libgomp.h         oacc-parallel.c  splay-tree.h
config            libgomp.map      oacc-plugin.c    target.c
config.h.in       libgomp-plugin.c oacc-plugin.h    task.c
configure         libgomp-plugin.h omp.h.in          taskloop.c
configure.ac       libgomp.spec.in  omp_lib.f90.in   team.c
configure.tgt     libgomp.texi     omp_lib.h.in     testsuite
critical.c        lock.c           openacc.f90      work.c
env.c             loop.c           openacc.h
error.c           loop_ull.c       openacc_lib.h
fortran.c         Makefile.am      ordered.c
```

Easy to spot at least one file per #pragma construct

#pragma omp parallel

Data scope attributes

```
void ex2 ()
{
    int id;
    int base = 5;
    int slaves = 0;

    #pragma omp parallel
    {
        id = omp_get_thread_num();
        slaves = omp_get_num_threads() - 1;

        if (id == 0)
            printf ("Master: base = %d.", base*slaves);
        else
            printf ("Slave: base = %d.", base*id);
    }
}
```

Call runtime to get thread ID:
Every thread sees a different value



Master and slave threads
access the same variable **a**



A slightly more complex example

#pragma omp parallel

Data scope attributes

```
void ex2 ()  
{  
    int id;  
    int base = 5;  
    int slaves = 0;
```

```
#pragma omp parallel
```

```
{
```

```
    id = omp_get_thread_num();  
    slaves = omp_get_num_procs() - 1;
```

```
    if (id == 0)
```

```
        printf ("Master: base = %d.", base*slaves);
```

```
    else
```

```
        printf ("Slave: base = %d.", base*id);
```

```
}
```

```
}
```

Call runtime to get thread ID:
Every thread sees a different value

How to inform the compiler
about these different
behaviors?

Master and slave threads
access the same variable **a**

A slightly more complex example

#pragma omp parallel

Data scope attributes

```
void ex2 ()
```

```
{
```

```
    int id;
```

```
    int base = 5;
```

```
    int slaves = 0;
```

Insert code to retrieve the address of the shared object from within each parallel thread



```
#pragma omp parallel shared (base, slaves) private (id)
```

```
{
```

```
    id = omp_get_thread_num();
```

```
    slaves = omp_get_num_threads() - 1;
```

Allow symbol privatization:
Each thread contains a private copy of this variable



```
    if (id == 0)
```

```
        printf ("Master: base = %d.", base*slaves);
```

```
    else
```

```
        printf ("Slave: base = %d.", base*id);
```

```
}
```

A slightly more complex example

```
}
```


More data sharing clauses

- **firstprivate**
 - copyin, private storage
- **lastprivate**
 - Copyout, private storage
- **Exercise** – See how the generated code differs in the uses of a **shared** and a **firstprivate** variable

OpenMP lowering

Stampa l'intermedio dopo il passo di
LOWERING di OpenMP



`gcc -fopenmp -fdump-tree-omplower main.c`

First of two stages of OpenMP processing:

1. Lowering
2. Expansion

Introduced new external node (ex2._omp_fn.1/6).

```
ex2 ()
{
  int id;
  int base;
  int slaves;

  base = 5;
  slaves = 0;
  {
    .omp_data_o.2.base = base;
    .omp_data_o.2.slaves = slaves;
    #pragma omp parallel private(id) shared(slaves) firstprivate(base) [child fn: ex2._omp_fn.1 (.omp_data_o.2)]
    {
      .omp_data_i = (struct .omp_data_s.1 & restrict) &.omp_data_o.2;
      base = .omp_data_i->base;
      id = omp_get_thread_num ();
      D.6171 = omp_get_num_threads ();
      D.6189 = D.6171 + -1;
      .omp_data_i->slaves = D.6189;
      if (id == 0) goto <D.6186>; else goto <D.6187>;
      <D.6186>:
      D.6190 = .omp_data_i->slaves;
      D.6174 = base * D.6190;
      printf ("Master: base = %d.", D.6174);
      goto <D.6188>;
      <D.6187>:
      D.6176 = base * id;
      printf ("Slave: base = %d.", D.6176);
      <D.6188>:
      #pragma omp return
    }
    slaves = .omp_data_o.2.slaves;
    .omp_data_o.2 = {CLOBBER};
  }
}
```

Operates in high
GIMPLE form

```
hello.c
hello.c.001t.tu
hello.c.002t.class
hello.c.003t.original
hello.c.004t.gimple
hello.c.006t.omplower
hello.c.007t.lower
hello.c.010t.eh
hello.c.011t.cfg
hello.c.012t.ompxd
```

`ls -la`

`hello.c.006t.omplower`

C Source Code

Parser

AST

Gimplifier

GIMPLE

Linearizer

Lower

CFG Generator

CFG

RTL Generator

RTL expand

OpenMP lowering

Introduced new external node (ex2. omp fn.1/6).

```
ex2 ()
{
  int id;
  int base;
  int slaves;

  base = 5;
  slaves = 0;

  {
    .omp_data_o.2.base = base;
    .omp_data_o.2.slaves = slaves;
    #pragma omp parallel private(id) shared(slaves) firstprivate(base) child fn: ex2. omp fn.1 (.omp_data_o.2)
    {
      .omp_data_i = (struct .omp_data_s.1 & restrict) &.omp_data_o.2;
      base = .omp_data_i->base;
      id = omp_get_thread_num ();
      D.6171 = omp_get_num_threads ();
      D.6189 = D.6171 + -1;
      .omp_data_i->slaves = D.6189;
      if (id == 0) goto <D.6186>; else goto <D.6187>;
      <D.6186>:
      D.6190 = .omp_data_i->slaves;
      D.6174 = base * D.6190;
      printf ("Master: base = %d.", D.6174);
      goto <D.6188>;
      <D.6187>:
      D.6176 = base * id;
      printf ("Slave: base = %d.", D.6176);
      <D.6188>:
      #pragma omp return
    }
    slaves = .omp_data_o.2.slaves;
    .omp_data_o.2 = {CLOBBER};
  }
}
```

In this stage we create a new function node...

...and implement **data marshalling** for sharing data among parallel threads

(but outlining doesn't happen just yet)

OpenMP lowering

Introduced new external node (ex2._omp_fn.1/6).

```
ex2 ()
```

```
{  
  int id;  
  int base;  
  int slaves;
```

```
  base = 5;  
  slaves = 0;
```

marshalling IN/OUT (omp_data_t omp_data_o;)

```
  .omp_data_o.2.base = base;  
  .omp_data_o.2.slaves = slaves;
```

```
  #pragma omp parallel private(id) shared(slaves) firstprivate(base) [child fn: ex2._omp_fn.1 (.omp_data_o.2)]
```

```
  {  
    .omp_data_i = (struct .omp_data_s.1 & restrict) &.omp_data_o.2;  
    base = .omp_data_i->base;
```

```
    id = omp_get_thread_num ();  
    D.6171 = omp_get_num_threads ();  
    D.6189 = D.6171 + -1;
```

```
    .omp_data_i->slaves = D.6189;
```

```
    if (id == 0) goto <D.6186>; else goto <D.6187>;  
    <D.6186>:
```

```
    D.6190 = .omp_data_i->slaves;
```

```
    D.6174 = base * D.6190;  
    printf ("Master: base = %d.", D.6174);  
    goto <D.6188>;
```

```
    <D.6187>:
```

```
    D.6176 = base * id;  
    printf ("Slave: base = %d.", D.6176);  
    <D.6188>:
```

```
    #pragma omp return
```

```
  }
```

```
  slaves = .omp_data_o.2.slaves;
```

```
  .omp_data_o.2 = {CLOBBER};
```

```
}
```

typedef struct

```
{  
  int base;  
  int slaves;  
} omp_data_t
```

...and implement **data marshalling** for sharing data among parallel threads

Replace uses (omp_data_t omp_data_i;)

hello.c.006t.omplower

OpenMP expansion

Stampa l'intermedio dopo il passo di
EXPANSION di OpenMP



```
ex2._omp_fn.1 (struct .omp_data_s.1 & restrict .omp_data_i)
```

```
{  
  int slaves [value-expr: .omp_data_i->slaves]  
  int D.6198;  
  int D.6197;  
  int D.6196;  
  int D.6195;  
  int D.6194;  
  int id;  
  int base;
```

```
<bb 3> [0.00%]:  
base = .omp_data_i->base;  
id = omp_get_thread_num ();  
D.6194 = omp_get_num_threads ();  
D.6195 = D.6194 + -1;  
.omp_data_i->slaves = D.6195;  
if (id == 0)  
  goto <bb 4>; [0.00%]  
else  
  goto <bb 5>; [0.00%]
```

```
<bb 6> [0.00%]:  
return;  
  
<bb 5> [0.00%]:  
D.6196 = base * id;  
printf ("Slave: base = %d.", D.6196);  
goto <bb 6>; [0.00%]
```

```
<bb 4> [0.00%]:  
D.6197 = .omp_data_i->slaves;  
D.6198 = base * D.6197;  
printf ("Master: base = %d.", D.6198);  
goto <bb 6>; [0.00%]
```

```
}
```

gcc -fopenmp -fdump-tree-ompexp main.c

Second of two stages of OpenMP processing:

1. Lowering
2. Expansion


Operates in low
GIMPLE form

```
hello.c  
hello.c.001t.tu  
hello.c.002t.class  
hello.c.003t.original  
hello.c.004t.gimple  
hello.c.006t.omplower  
hello.c.007t.lower  
hello.c.010t.eh  
hello.c.011t.cfg  
hello.c.012t.ompexp
```


ls -la

hello.c.012t.ompexp


C Source Code



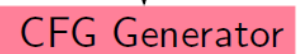
AST




GIMPLE



Lower



CFG



RTL expand

OpenMP expansion

Stampa l'intermedio dopo il passo di
EXPANSION di OpenMP



```
ex2._omp_fn.1 (struct .omp_data_s.1 & restrict .omp_data_i)
```

```
{  
    int slaves [value-expr: .omp_data_i->slaves]  
    int D.6198;  
    int D.6197;  
    int D.6196;  
    int D.6195;  
    int D.6194;  
    int id;  
    int base;
```

```
<bb 3> [0.00%]:  
base = .omp_data_i->base;  
id = omp_get_thread_num ();  
D.6194 = omp_get_num_threads ();  
D.6195 = D.6194 + -1;  
.omp_data_i->slaves = D.6195;  
if (id == 0)  
    goto <bb 4>; [0.00%]  
else  
    goto <bb 5>; [0.00%]
```

```
<bb 6> [0.00%]:  
return;
```

```
<bb 5> [0.00%]:  
D.6196 = base * id;  
printf ("Slave: base = %d.", D.6196);  
goto <bb 6>; [0.00%]
```

```
<bb 4> [0.00%]:  
D.6197 = .omp_data_i->slaves;  
D.6198 = base * D.6197;  
printf ("Master: base = %d.", D.6198);  
goto <bb 6>; [0.00%]
```

```
}
```

gcc -fopenmp -fdump-tree-ompexp main.c
-fdump-tree-all per stampare la IR dopo ogni
passo di ottimizzazione del middle-end

Second of two stages of OpenMP processing:


1. Lowering
2. Expansion

Operates in low
GIMPLE form

```
hello.c  
hello.c.001t.tu  
hello.c.002t.class  
hello.c.003t.original  
hello.c.004t.gimple  
hello.c.006t.omplower  
hello.c.007t.lower  
hello.c.010t.eh  
hello.c.011t.cfg  
hello.c.012t.ompexp
```

ls -la

hello.c.012t.ompexp



OpenMP expansion

```
ex2._omp_fn.1 (struct .omp_data_s.1 & restrict .omp_data_i)
```

```
{  
    int slaves [value-expr: .omp_data_i->slaves];  
    int D.6198;  
    int D.6197;  
    int D.6196;  
    int D.6195;  
    int D.6194;  
    int id;  
    int base;
```

```
<bb 3> [0.00%]:
```

```
base = .omp_data_i->base;
```

```
id = omp_get_thread_num();
```

```
D.6194 = omp_get_num_threads();
```

```
D.6195 = D.6194 + -1;
```

```
.omp_data_i->slaves = D.6195;
```

```
if (id == 0)
```

```
    goto <bb 4>; [0.00%]
```

```
else
```

```
    goto <bb 5>; [0.00%]
```

```
<bb 6> [0.00%]:
```

```
return;
```

```
<bb 5> [0.00%]:
```

```
D.6196 = base * id;
```

```
printf ("Slave: base = %d.", D.6196);
```

```
goto <bb 6>; [0.00%]
```

```
<bb 4> [0.00%]:
```

```
D.6197 = .omp_data_i->slaves;
```

```
D.6198 = base * D.6197;
```

```
printf ("Master: base = %d.", D.6198);
```

```
goto <bb 6>; [0.00%]
```

marshalling IN/OUT (omp_data_t omp_data_o;)

```
ex2 ()
```

```
{
```

```
    int slaves;
```

```
    int base;
```

```
    int id;
```

```
    struct .omp_data_s.1 .omp_data_o.2;
```

```
    int D.6190;
```

```
    int D.6189;
```

```
<bb 2> [0.00%]:
```

```
base = 5;
```

```
slaves = 0;
```

```
.omp_data_o.2.base = base;
```

```
.omp_data_o.2.slaves = slaves;
```

```
__builtin_GOMP_parallel (ex2._omp_fn.1, &.omp_data_o.2, 0, 0);
```

```
slaves = .omp_data_o.2.slaves;
```

```
.omp_data_o.2 = {CLOBBER};
```

```
return;
```

```
}
```

Call runtime by passing a
pointer to the outlined function
and the marshalled data

Replace uses (omp_data_t omp_data_i;)

hello.c.012t.ompexp

Sharing work among threads

The **for** directive

- The **parallel** pragma instructs every thread to execute all of the code inside the block
- If we encounter a **for** loop that we want to divide among threads, we use the **for** pragma

```
#pragma omp for
```


#pragma omp for

#pragma omp for can be placed everywhere inside a **parallel** construct, or combined with it, as in the example

The code of the **for** loop is moved inside the outlined function.

```
void ex3 (int *a)
{
    int i;

    #pragma omp parallel for
    for (i=0; i<10; i++)
        a[i] = i;
}
```


```
void ex3 (...)
{
    omp_parallel_start(&parfun, ...);
    parfun();
    omp_parallel_end();
}

int parfun(...)
{
    int LB = ...;
    int UB = ...;

    for (i=LB; i<UB; i++)
        a[i] = i;
}
```

#pragma omp for

Every thread works on a different subset of the iteration space..



```
void ex3 (int *a)
{
    int i;

    #pragma omp parallel for
    for (i=0; i<10; i++)
        a[i] = i;
}
```


```
void ex3(...)
{
    omp_parallel_start(&parfun, ...);
    parfun();
    omp_parallel_end();
}

int parfun(...)
{
    int LB = ...;
    int UB = ...;

    for (i=LB; i<UB; i++)
        a[i] = i;
}
```

#pragma omp for

..since lower and upper bounds (LB, UB) are computed locally as a function of the thread ID



```
void ex3 (int *a)
{
    int i;

    #pragma omp parallel for
    for (i=0; i<10; i++)
        a[i] = i;
}
```

```
void ex3 (...)
{
    omp_parallel_start(&parfun, ...);
    parfun();
    omp_parallel_end();
}

int parfun(...)
{
    int LB = ...;
    int UB = ...;

    for (i=LB; i<UB; i++)
        a[i] = i;
}
```

The **schedule** clause

Static Loop Partitioning

```
#pragma omp for schedule(static)
{
    for (i=0; i<12; i++)
        a[i] = i;
}
```

Useful for:


- *Simple, regular loops*
- *Iterations with equal duration*

Es. 12 iterations (N), 4 threads (Nthr)

DATA CHUNK

$$C = \text{ceil} \left(\frac{N}{N_{\text{thr}}} \right)$$

3 iterations thread



Thread ID (TID)		0	1	2	3
LOWER BOUND	$LB = C * TID$	0	3	6	9
UPPER BOUND	$UB = \min \{ [C * (TID + 1)], N \}$	3	6	9	12

The **schedule** clause

Static Loop Partitioning

```
#pragma omp for schedule(static)
{
    for (i=0; i<12; i++)
        a[i] = i;
}
```

Useful for:


- *Simple, regular loops*
- *Iterations with equal duration*

Es. 12 iterations (N), 4 threads (Nthr)

DATA CHUNK

$$C = \text{ceil} \left(\frac{N}{N_{\text{thr}}} \right)$$

3 iterations thread



Thread ID (TID)		0	1	2	3
LOWER BOUND	$LB = C * TID$	0	3	6	9
UPPER BOUND	$UB = \min \{ [C * (TID + 1)], N \}$	3	6	9	12

Static scheduling clauses

- **EXERCISE** – Generate and analyze the IR dumps for the OpenMP lowering and expansion passes for the `#pragma omp for` directive, playing with static chunking
- `schedule (static, N)`

The **schedule** clause

Static Loop Partitioning

Es. 12 iterations (N), 4 threads (Nthr)

```
#pragma omp for schedule(static)
{
    for (i=0; i<12; i++)
        a[i] = i;
}
```

Useful for:

- Simple, regular loop
- Iterations with equal work

DATA CHUNK

$$e_{il} \left(\frac{N}{N_{thr}} \right)$$

**3
iterations
thread**

What happens with static scheduling when iterations have different duration?

LOWER BO

UPPER BOUND

 ~~$\{C^*(TID + 1), N\}$~~

	0	1	2	3	
0					
3					

The **schedule** clause

Static Loop Partitioning

```
#pragma omp for schedule(static)
```

```
{  
  for (i=0; i<12; i++)
```

```
{
```

```
    int start = rand();
```

```
    int count = 0;
```



```
    while (start++ < 256)  
        count++;
```

```
    a[count] = foo();
```

```
}
```

```
}
```

A variable amount of work in
each iteration



UNBALANCED
workloads


The **schedule** clause

Dynamic Loop Partitioning

```
#pragma omp for schedule(dynamic)
{
    for (i=0; i<12; i++)
    {
        int start = rand();
        int count = 0;

        while (start++ < 256)
            count++;

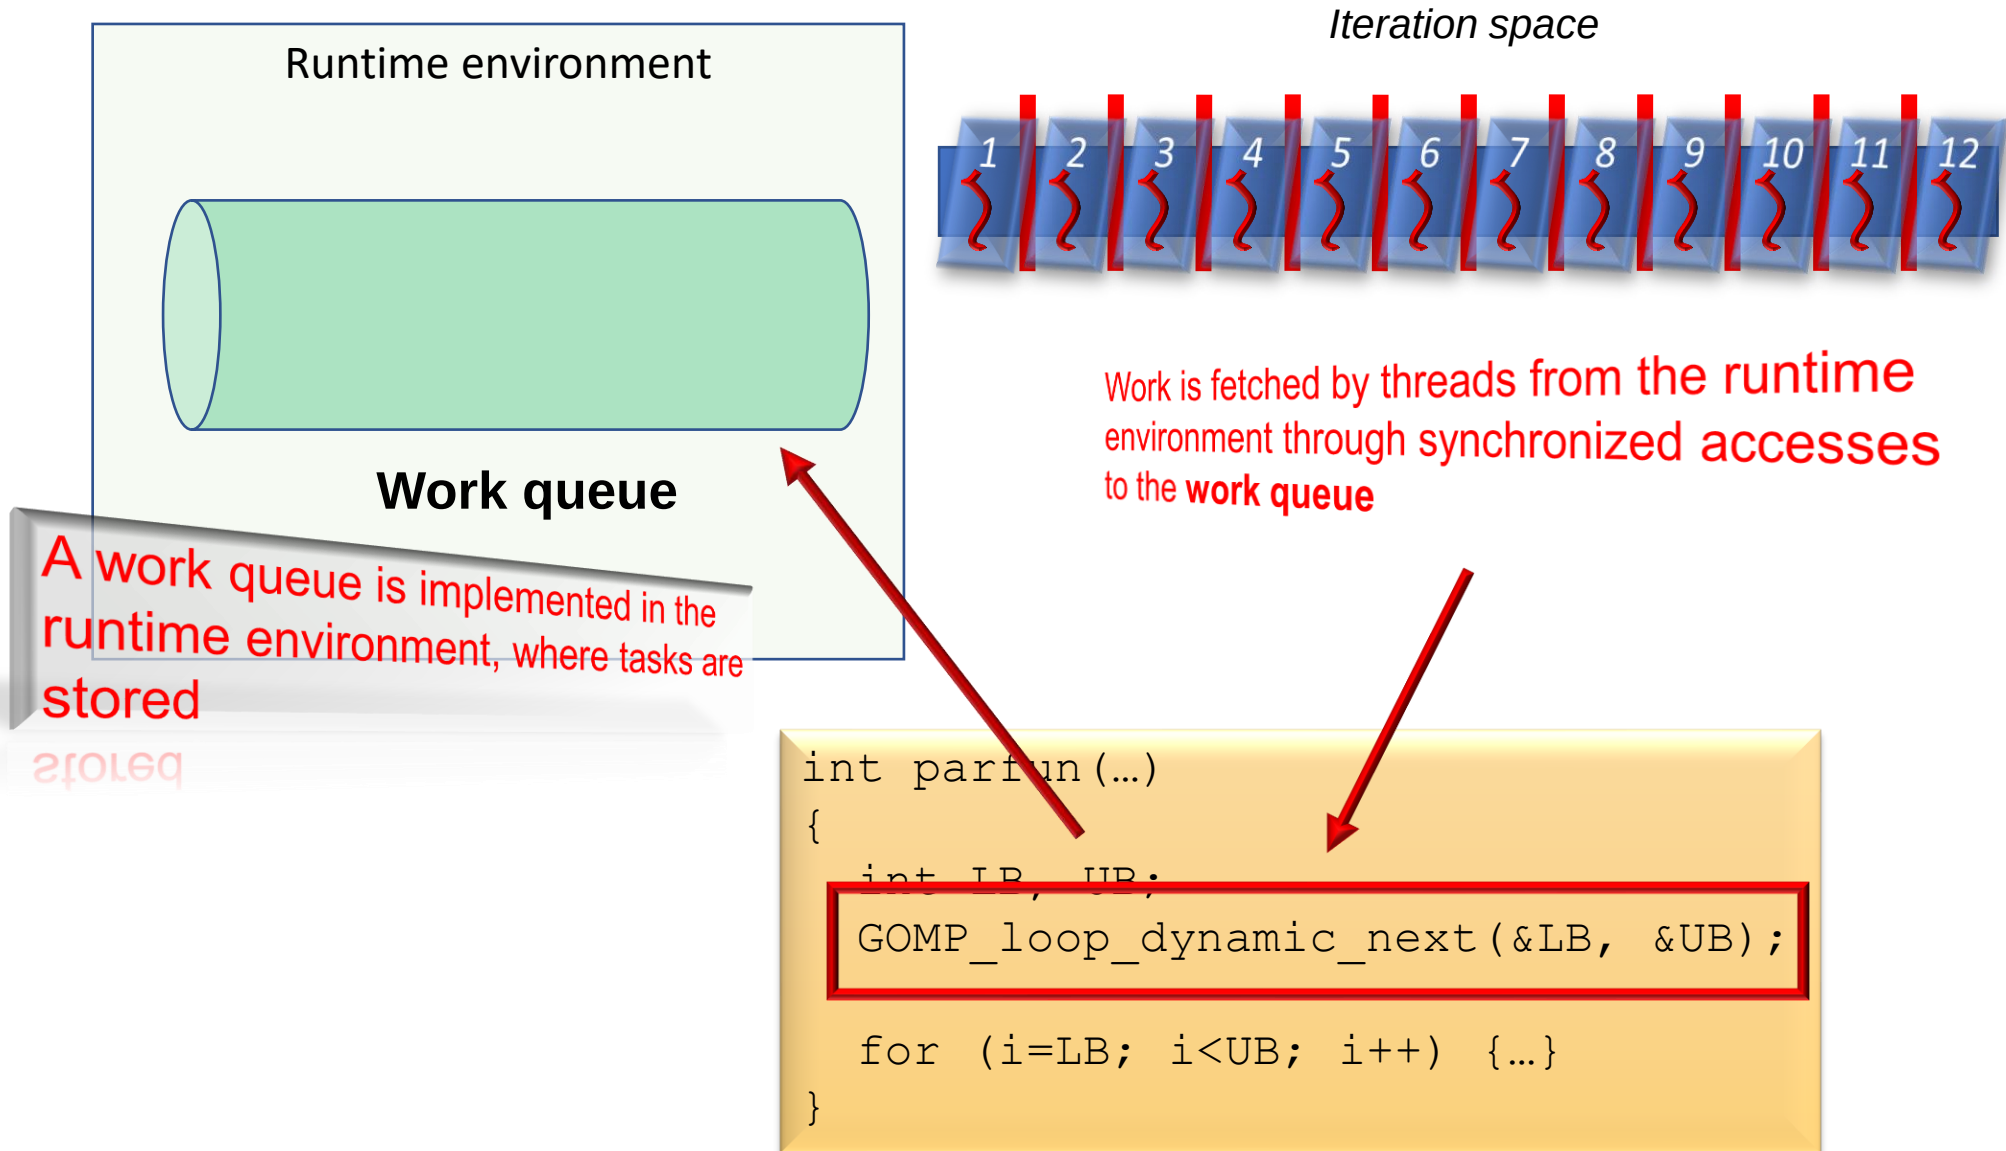
        a[count] = foo();
    }
}
```



A thread is generated for every
single iteration

The **schedule** clause

Dynamic Loop Partitioning






The **schedule** clause

Dynamic Loop Partitioning

Remember results with
static scheduling..

Iteration space



Speedup!

BALANCED
workloads

Parallelization granularity

Iteration chunking

- Fine-grain Parallelism

- Best opportunities for load balancing, but..
- Small amounts of computational work between parallelism computation stages
- Low computation to parallelization ratio → High parallelization overhead



- Coarse-grain Parallelism

- Harder to load balance efficiently, but..
- Large amounts of computational work between parallelism computation stages
- High computation to parallelization ratio → Low parallelization overhead



The **schedule** clause


Dynamic Loop Partitioning

```
#pragma omp for schedule(dynamic, 1)
{
    for (i=0; i<12; i++)
    {
        int start = rand();
        int count = 0;

        while (start++ < 256)
            count++;



        a[count] = foo();
    }
}
```

Iteration space



The **schedule** clause

Dynamic Loop Partitioning



Runtime scheduling primitive is invoked at every iteration

If granularity of WORK is very small the overhead for fine chunking is significant




```
int parfun(...)  
{  
    int LB, UB;  
    GOMP_loop_dynamic_next(&LB, &UB);  
    for (i=LB; i<UB; i++) {WORK}  
}
```

The **schedule** clause

Dynamic Loop Partitioning

Remember results with
static scheduling..

Iteration space





The **schedule** clause

Dynamic Loop Partitioning

chunking overhead

Iteration space



The **schedule** clause

Dynamic Loop Partitioning

```
#pragma omp for schedule(dynamic, 2)
```

```
{
```

```
  for (i=0; i<12; i++)
```

```
  {
```

```
    int start = rand();
```

```
    int count = 0;
```

```
    while (start++ < 256)
```


```
      count++;
```

```
    a[count] = foo();
```

```
  }
```



```
}
```

Iteration space



The **schedule** clause

Dynamic Loop Partitioning



Runtime scheduling primitive is invoked every two iterations (half the times of previous case)

WORK is repeated twice among two scheduling events. Overhead gets amortized




```
int parfun(...)
{
    int LB, UB;
    GOMP_loop_dynamic_next(&LB, &UB);
    for (i=LB; i<UB; i++) {WORK}
}
```

The **schedule** clause

Dynamic Loop Partitioning

Remember results with
static scheduling..

Iteration space




The **schedule** clause

Dynamic Loop Partitioning

Iteration space

smaller chunking
overhead



More scheduling clauses

- **schedule (guided[, *chunk*])**
 - Threads dynamically grab blocks of iterations. The size of the block starts large and shrinks down to size “chunk” as the calculation proceeds.
- **schedule (runtime[, *chunk*])**
 - Schedule and chunk size taken from the OMP_SCHEDULE environment variable (or the runtime library ... for OpenMP 3.0)
- **EXERCISE** – Generate the IR dumps for the OpenMP lowering and expansion passes for different scheduling clauses


#pragma omp critical

- **Critical Section**: a portion of code that only one thread at a time may execute
- We denote a critical section by putting the pragma

```
#pragma omp critical
```

in front of a block of C code

π -finding code example




```
double area, pi, x;
int i, n;
#pragma omp parallel for private(x) \
                        shared(area)
{
    for (i=0; i<n; i++) {
        x = (i + 0.5)/n;
        area += 4.0/(1.0 + x*x);
    }
    pi = area/n;
}
```

Synchronize accesses to
shared variable **area** to
avoid inconsistent results


Race condition

- Ensure atomic updates of the shared variable **area** to avoid a **race condition** in which one process may “race ahead” of another and ignore changes




Race condition (Cont'd)

- Thread A reads "11.667" into a local register
- Thread B reads "11.667" into a local register
- Thread A updates area with "11.667+3.765"
- Thread B ignores write from thread A and updates area with "11.667 + 3.563"



π -finding code example

```
double area, pi, x;
int i, n;
#pragma omp parallel for private(x) shared(area)
{
    for (i=0; i<n; i++) {
        x = (i + 0.5)/n;
        #pragma omp critical
        area += 4.0 / (1.0 + x*x);
    }
}
pi = area/n;
```



#pragma omp critical protects the code within its scope by acquiring a lock before entering the critical section and releasing it after execution

Correctness and performance!

- As a matter of fact, using locks makes execution **sequential**
- To dim this effect we should try use **fine grained locking** (i.e. make critical sections as small as possible)
- A simple instruction to compute the value of area in the previous example is translated into many more simpler instructions within the compiler!
- The programmer is not aware of the real **granularity** of the critical section

```

_suif_start.omp_fn.0 (.omp_data_i)
{
  <bb 2>:
    D.2605 = .omp_data_i->n;
    D.2615 = __builtin_omp_get_num_threads ();
    D.2616 = __builtin_omp_get_thread_num ();

    ***
    D.2623 = MIN_EXPR <D.2622, D.2605>;
    if (D.2621 >= D.2623) goto <L4>; else goto <L2>;

  <L4>::
    return;

  <L2>::
    D.2586 = (double) i;
    D.2587 = D.2586 + 5.0e-1;
    D.2605 = .omp_data_i->n;
    D.2606 = D.2605;
    D.2588 = (double) D.2606;
    x = D.2587 / D.2588;
    __builtin_GOMP_critical_start ();
    D.2589 = x * x;
    D.2590 = D.2589 + 1.0e+0;
    D.2591 = 4.0e+0 / D.2590;
    D.2607 = .omp_data_i->area;
    D.2608 = D.2607;
    D.2609 = D.2591 + D.2608;
    .omp_data_i->area = D.2609;
    __builtin_GOMP_critical_end ();
    i = i + 1;
    D.2627 = i < D.2626;
    if (D.2627) goto <L2>; else goto <L4>;
}

```

execution sequential

This is a dump of the intermediate representation of the program within the compiler

```

_suif_start.omp_fn.0 (.omp_data_i)
{
  <bb 2>:
    D.2605 = .omp_data_i->n;
    D.2615 = __builtin_omp_get_num_threads ();
    D.2616 = __builtin_omp_get_thread_num ();

    ***
    D.2623 = MIN_EXPR <D.2622, D.2605>;
    if (D.2621 >= D.2623) goto <L4>; else goto <L2>;

```

```

<L4>::
  return;

```

```

<L2>::
  D.2586 = (double) i;
  D.2587 = D.2586 + 5.0e-1;
  D.2605 = .omp_data_i->n;
  D.2606 = D.2605;
  D.2588 = (double) D.2606;
  x = D.2587 / D.2588;
  __builtin_GOMP_critical_start ();
  D.2589 = x * x;
  D.2590 = D.2589 + 1.0e+0;
  D.2591 = 4.0e+0 / D.2590;
  D.2607 = .omp_data_i->area;
  D.2608 = D.2607;
  D.2609 = D.2591 + D.2608;
  .omp_data_i->area = D.2609;
  __builtin_GOMP_critical_end ();
  i = i + 1;
  D.2627 = i < D.2626;
  if (D.2627) goto <L2>; else goto <L4>;

```

execution sequential
grained locking (i.e.
e)

of area in the
/ more simpler

granularity of the

*This is how the compiler represents
the instruction*

*area += 4.0 / (1.0 + x*x);*

3,0-1


```
_suif_start.omp_fn.0 (.omp_data_i)  
{
```

```
<bb 2>:
```

```
D.2605 = .omp_data_i->n;
```

```
D.2615 = __builtin_omp_get_num_threads();
```

```
D.2616 = __builtin_omp_get_thread_num();
```

```
***  
D.2623 = MIN_EXPR <D.2622, D.2605>;
```

```
if (D.2621 >= D.2623) goto <L4>; else goto <L2>;
```

```
<L4>::
```

```
return;
```

```
<L2>::
```

```
D.2586 = (double) i;
```

```
D.2587 = D.2586 + 5.0e-1;
```

```
D.2605 = .omp_data_i->n;
```

```
D.2606 = D.2605;
```

```
D.2588 = (double) D.2606;
```

```
x = D.2587 / D.2588;
```

```
__builtin_GOMP_critical_start();
```

```
D.2589 = x * x;
```

```
D.2590 = D.2589 + 1.0e+0;
```

```
D.2591 = 4.0e+0 / D.2590;
```

```
D.2607 = .omp_data_i->area;
```

```
D.2608 = D.2607;
```

```
D.2608 = D.2591 + D.2608;
```

```
D.2609 = D.2591 * D.2609;
```

```
__builtin_GOMP_critical_end();
```

```
i = i + 1;
```

```
D.2627 = i < D.2626;
```

```
if (D.2627) goto <L2>; else goto <L4>;
```

**THIS IS DONE AT EVERY
LOOP ITERATION!**

call runtime to acquire lock

Lock-protected
operations (**critical
section**)

call runtime to release lock

coarse-grained locking (i.e. low
granularity)

fine-grained locking (i.e. high
granularity)

granularity of the
locking

3,0-1

π -finding code example

```
double area, pi, x;
int i, n;
#pragma omp parallel for \
    private(x) \
    shared(area)
{
    for (i=0; i<n; i++) {
```

```
        x = (i + 0.5) / n;
```

Parallel

```
        #pragma omp critical
```

```
        area += 4.0 / (1.0 + x*x);
```

Sequential


```
    }
```

```
}
```

```
pi = area/n;
```



Waiting for lock




Correctness and performance!

- A programming pattern such as `area += 4.0 / (1.0 + x*x) ;` in which we:
 - *Fetch the value of an operand*
 - *Add a value to it*
 - *Store the updated value*is called a **reduction**, and is commonly supported by parallel programming APIs
- OpenMP takes care of storing partial results in **private variables** and combining partial results after the loop

Correctness and performance!

```
double area, pi, x;
int i, n;
#pragma omp parallel for private(x) shared(area) reduction(+:area)
{
    for (i=0; i<n; i++) {
        x = (i + 0.5)/n;

        area += 4.0/(1.0 + x*x);
    }
}
pi = area/n;
```



The **reduction** clause instructs the compiler to create **private** copies of the **area** variable for every thread. At the end of the loop partial sums are combined on the shared **area** variable

ance!

```
_suif_start.omp_fn.0 (.omp_data_i)
{
  <bb 2>:
    area = 0.0;
    D.2605 = .omp_data_i->n;
    D.2615 = __builtin_omp_get_num_threads ();
    D.2616 = __builtin_omp_get_thread_num ();

    ***
    D.2623 = MIN_EXPR <D.2622, D.2605>;
    if (D.2621 >= D.2623) goto <L4>; else goto <L2>;
```

```
<L4>::
  __builtin_GOMP_atomic_start ();
  D.2607 = &.omp_data_i->area;
  D.2608 = *D.2607;
  D.2609 = D.2608 + area;
  *D.2607 = D.2609;
  __builtin_GOMP_atomic_end ();
  return;
```

```
<L2>::
  D.2586 = (double) i;
  D.2587 = D.2586 + 5.0e-1;
  D.2605 = .omp_data_i->n;
  D.2606 = D.2605;
  D.2588 = (double) D.2606;
  x = D.2587 / D.2588;
  D.2589 = x * x;
  D.2590 = D.2589 + 1.0e+0;
  D.2591 = 4.0e+0 / D.2590;
  area = D.2591 + area;
  area = i + 1;
  i = i + 1;
  D.2627 = i < D.2626;
  if (D.2627) goto <L2>; else goto <L4>;
```

Shared variable is only updated at the end of the loop, when partial sums are computed

shared(area) reduction(+:area)

Each thread computes partial sums on private copies of the reduction variable

The r
the a
comb

er to create **private** copies of
nd of the loop partial sums are

ance!

double
int

#pragma
{

for

a

}

}

pi

```
_suif_start.omp_fn.0 (.omp_data_i)
{
  <bb 2>:
    area = 0.0;
    D.2605 = .omp_data_i->n;
    D.2615 = __builtin_omp_get_num_threads ();
    D.2616 = __builtin_omp_get_thread_num ();

    ***
    D.2623 = MIN_EXPR <D.2622, D.2605>;
    if (D.2621 >= D.2623) goto <L4>; else goto <L2>;

    <L4>::
      __builtin_GOMP_atomic_start ();
      D.2607 = &.omp_data_i->area;
      D.2608 = *D.2607;
      D.2609 = D.2608 + area;
      *D.2607 = D.2609;
      __builtin_GOMP_atomic_end ();
      return;


    <L2>::
      D.2586 = (double) i;
      D.2587 = D.2586 + 5.0e-1;
      D.2605 = .omp_data_i->n;
      D.2606 = D.2605;
      D.2588 = (double) D.2606;
      x = D.2587 / D.2588;
      D.2589 = x * x;
      D.2590 = D.2589 + 1.0e+0;
      D.2591 = 4.0e+0 / D.2590;
      area = D.2591 + area;
      area = i + 1;
      i = i + 1;
      D.2627 = i < D.2626;
      if (D.2627) goto <L2>; else goto <L4>;
}
```

```
__builtin_GOMP_atomic_start ();
D.2607 = &.omp_data_i->area;
D.2608 = *D.2607;
D.2609 = D.2608 + area;
*D.2607 = D.2609;
__builtin_GOMP_atomic_end ();
return;
```

```
<L2>::
D.2586 = (double) i;
D.2587 = D.2586 + 5.0e-1;
D.2605 = .omp_data_i->n;
D.2606 = D.2605;
D.2588 = (double) D.2606;
x = D.2587 / D.2588;
D.2589 = x * x;
D.2590 = D.2589 + 1.0e+0;
D.2591 = 4.0e+0 / D.2590;
area = D.2591 + area;
area = i + 1;
i = i + 1;
D.2627 = i < D.2626;
if (D.2627) goto <L2>; else goto <L4>;
```

The r
the a
comb

shared area)



er to create private copies of
nd of the loop partial sums are